

# Supplementary Material for “Direct Quantized Training of Language Models with Stochastic Rounding”

Kaiyan Zhao<sup>1</sup>

KAIYAN1006@LOGOS.T.U-TOKYO.AC.JP

Tsuguchika Tabaru<sup>2</sup>

TABARU@FUJITSU.COM

Kenichi Kobayashi<sup>2</sup>

KENICHI@FUJITSU.COM

Takumi Honda<sup>2</sup>

HONDA.TAKUMI@FUJITSU.COM

Masafumi Yamazaki<sup>2</sup>

M.YAMAZAKI@FUJITSU.COM

Yoshimasa Tsuruoka<sup>1</sup>

TSURUOKA@LOGOS.T.U-TOKYO.AC.JP

<sup>1</sup> The University of Tokyo, Tokyo, Japan

<sup>2</sup> Fujitsu Limited, Kawasaki, Japan

1 In this supplementary material, we provide additional details to complement the main  
2 paper. Specifically, we present:

- 3 • Additional implementation details;
- 4 • Additional experimentl results including non-logarithmic comparison;
- 5 • Explanation of our models with ternary inference;
- 6 • A proof sketch for stochastic rounding’s convergence guarantee in DQT

## 7 1. Additional Implementation Details

### 8 1.1. Model configuration

9 The detailed configurations for models of different sizes are provided in Table 1. Note that  
10 the batch size varies across model sizes during training, but no gradient accumulation is  
11 used in any experiment. For each model, the learning rate is selected via grid search over  
12 the set {1e-5, 1e-4, 5e-4, 1e-3} using our development set. Regarding the tokenizer, we  
13 adopt a publicly released pre-trained one<sup>1</sup> without further updates during training. We fix  
14 the random seed to 42 for reproducibility.

### 15 1.2. Dataset preprocessing

16 The maximum length of training data for both datasets is set to 512. Texts longer than  
17 512 tokens are split into separate chunks, while shorter texts are padded accordingly. For  
18 Wikipedia, this preprocessing results in approximately 14 million sentences derived from  
19 the original 6.4 million examples, while for FineWeb, the dataset yields around 33 million  
20 sentences. The test set of WikiText-2 is from <sup>2</sup>.

---

1. [https://huggingface.co/1bitLLM/bitnet\\_b1\\_58-large](https://huggingface.co/1bitLLM/bitnet_b1_58-large)

2. <https://huggingface.co/datasets/Salesforce/wikitext/viewer/wikitext-2-v1/test>

Params	hidden_size	intermediate_size	num_hidden_layers	num_attention_heads	batch_size
130M	768	2048	12	12	64
320M	1024	2048	24	16	32
1B	2048	3072	24	32	16

Table 1: Configuration of different models sizes.

Model size	FP32	BF16	BF16+Adafactor	FP8	FP8+Adafactor
130M	69327	54675	53827	39276	38315
1B	76533	58345	53723	40945	37669

Table 2: Actual GPU memory usage (in MB) of models with different sizes on a single GH200.

21 **1.3. Low-precision environments**

22 For FP8 simulation, we choose to use MS-AMP because the GH200 Superchips are equipped  
 23 with ARM-based CPUs, which are not fully supported by the transformer-engine library  
 24 currently. As for the reason for using low-precision simulation, hardware and software  
 25 constraints, such as the inability to modify low-level PyTorch kernels or implement true  
 26 integer-based weight updates, make it infeasible to perform actual low-bit computation in  
 27 our environment. These limitations are typical in academic settings, where direct access to  
 28 low-level hardware accelerators is restricted. As an alternative, we focus on evaluating the  
 29 practical effectiveness of our method under memory-constrained conditions, simulating the  
 30 scenarios where computational resources are limited.

31 **1.4. Actual Memory Usage**

32 We present the actual memory usage of models with different sizes in Table 2. The values  
 33 reflect usage on a single GH200 GPU, which has 97,871 MB of available memory. We can  
 34 also observe exactly how much memory is saved after applying the Adafactor optimizer.

35 **1.5. Weight Update Frequency**

36 For BitNet models, we compare quantized weight matrices at adjacent training steps by  
 37 iterating through all corresponding rows and columns to identify whether each element is  
 38 updated or not. For DQT models, we can directly compare the weight matrices before  
 39 and after stochastic rounding. The update percentage is computed as the ratio of changed  
 40 weight elements to the total number of elements.

41 **2. Additional Experimental Results**

42 We provide a more precise and clear comparison in Figure 1 and Figure 2.

43 **3. Ternary Inference**

44 Since the straight-through estimator is not employed in our proposed DQT, both the forward  
 45 and backward operate directly on  $n$ -bit weight matrices, which means that our DQT models

SUPPLEMENTARY MATERIAL




Figure 1: Non-logarithmic training loss comparison of DQT 8 bits and BitNet b1.58 in 1B sizes. Our DQT 8 bits performs slightly better than BitNet b1.58.




Figure 2: Training loss comparison of DQT 8 bits and DQT 8 bits that utilizes ternary inference. DQT achieves ternary inference with minimal degradation.

46 trained with larger bit-width are not inherently ternary during inference. To enable a fair  
 47 comparison with BitNet b1.58, we adapt our models to perform ternary inference. This is  
 48 achieved by using ternary weights during the forward pass while maintaining  $n$ -bit weights  
 49 for the backward pass, updated through the straight-through estimator. The straight-  
 50 through estimator is employed only to enable ternary inference in  $n$ -bit DQT models as a  
 51 variant of our models. We additionally present the training loss for DQT 8 bit and DQT 8  
 52 bit with ternary inference in the supplementary material, Figure 2. Note that when ternary  
 53 inference is applied, the memory usage is the same as BitNet b1.58 during inference.

#### 54 4. A Proof Sketch for Convergence Guarantee in DQT

55 We consider the simplified optimization problem:  $\min L(\theta)$ , where  $L(\cdot)$  is a smooth and  
56 possibly non-convex loss function and  $\theta$  stands for the model parameter.

57 We use a stochastic gradient method with quantized updates:

$$\theta_{t+1} = SR(\theta - \eta \nabla L(\theta_t)) = \theta_t - \eta \nabla L(\theta_t) + \epsilon_t, \quad (1)$$

58 where  $SR(\cdot)$  denotes the stochastic rounding quantizer (core approach of DQT),  $\eta$  is the  
59 learning rate and  $\epsilon_t$  is the quantization error (noise) which can be represented as:

$$\epsilon_t = SR(\theta - \eta \nabla L(\theta_t)) - (\theta - \eta \nabla L(\theta_t)). \quad (2)$$

60 It is easy to find that in Equation 1, the latter is in the form of SGD optimization with  
61 noise. Let  $x = \theta - \eta \nabla L(\theta_t)$ , then  $\epsilon_t$  can be written into:

$$\epsilon_t = SR(x) - x, \quad (3)$$

62 where

$$SR(x) = \begin{cases} \lfloor x \rfloor, & \text{with } p = \lceil x \rceil - x \\ \lceil x \rceil, & \text{otherwise} \end{cases}. \quad (4)$$

63 Now we can calculate the expectation:

$$\mathbb{E}(SR(x)) = (\lceil x \rceil - x) \cdot \lfloor x \rfloor + (x - \lfloor x \rfloor) \cdot \lceil x \rceil = x. \quad (5)$$

64 Then it would be obvious that

$$\mathbb{E}(\epsilon_t) = \mathbb{E}(SR(x)) - \mathbb{E}(x) = 0. \quad (6)$$

65 This draws the important conclusion that the noise is **zero-mean and unbiased**.

66 Next, we continue to calculate the variance of  $\epsilon_t$ . Assume we quantize  $x$  with step size  
67  $\Delta$ . Define:  $q_{low} = \lfloor x/\Delta \rfloor \cdot \Delta$ , and  $q_{high} = q_{low} + \Delta$ . With probability  $\alpha$ ,  $SR(x) = q_{high}$   
68 and with probability  $1 - \alpha$ ,  $SR(x) = q_{low}$ . Then following Equation 3, the noise  $\epsilon_t$  can be  
69 represented as:

$$\epsilon_t = \begin{cases} q_{low} - x = -\alpha\Delta, & \text{with } p = 1 - \alpha \\ q_{high} - x = (1 - \alpha)\Delta, & \text{with } p = \alpha \end{cases}. \quad (7)$$

70 The variance of  $\epsilon_t$  can be represented using the following equation:

$$\mathbb{V}(\epsilon_t) = (1 - \alpha)(\alpha\Delta)^2 + \alpha[(1 - \alpha)\Delta]^2 = \alpha(1 - \alpha)\Delta^2. \quad (8)$$

71 Thus, the variance of  $\epsilon_t$  is bounded:

$$\mathbb{V}(\epsilon_t) \leq \frac{1}{4}\Delta^2, \quad (9)$$

72 the maximum is achieved when  $\alpha = \frac{1}{2}$ .

73 In conclusion, for an SGD optimization problem with noise, if the noise is **zero-mean**  
74 and its variance is bounded, the convergence can be guaranteed.

SUPPLEMENTARY MATERIAL




Figure 3: An example of our training process using 8-bit quantization.